

Machine Learning Project

Joshua Peterson

10/4/2016

Background

In this project I analyzed data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. This data may have come from IoT, fitness devices such as Jawbone Up, Nike FuelBand and the Fitbit. Each participant was asked to perform barbell lifts correctly and incorrectly in 5 different ways.

The intent of the project is to predict the manner in which each of the participants performed each of the performed exercises. This variable is the “classe” variable in the training set.

The training data for the project is available at:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

The test data for the project is available at:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

Study Design

I will be following the standard Prediction Design Framework

1. Define error rate
2. Splitting the data in to Training, Testing and Validation
3. Choose features of the Training data set using cross-validation
4. Choose prediction function of the Training data using cross-validation
5. If there is no validation apply 1x to test set
6. If there is validation apply to test set and refine and apply 1x to validation

It appears that we have a relatively large sample size therefore I would like to target the following parameters:

1. 60% training
2. 20% test
3. 20% validation

Loading the exercise data.

Data Splitting

In this data splitting action, I am splitting the data into 60% training set and 40% testing set.

```
library(caret); data(train_data)
```

```
## Error in read.table(file = file, header = header, sep = sep, quote = quote, : 'file' must be a character
```

```
inTrain <- createDataPartition(y=train_data$classe,p=0.60, list=FALSE)
training <- train_data[inTrain,]
testing <- train_data[-inTrain,]
dim(training); dim(testing)
```

```
## [1] 11776 160
```

```
## [1] 7846 160
```

Cleaning the data

Because there are some non-zero values within the data set, we must clean the data set to extract those values.

```
train_data_NZV<- nearZeroVar(training, saveMetrics = TRUE)
```

```
## [1] 11776 100
```

Next, because there are data points within the first column that may interfere with my algorithms, I will remove the first column from the data set

```
training<- training[c(-1)]
```

Next, I will eliminate those variables with an excessive amount of NAs values.

```
final_training <- training
for(i in 1:length(training)) {
  if( sum( is.na( training[, i] ) ) /nrow(training) >= .6 ) {
    for(j in 1:length(final_training)) {
      if( length( grep(names(training[i]), names(final_training)[j]) ) ==1 ) {
        }
      }
    }
  }
}
```

```
dim(final_training)
```

```
## [1] 11776 58
```

```
training<- final_training
rm(final_training)
```

I will now perform the same data cleansing process for the testing data that I performed for the training data.

```
clean_training<- colnames(training)
clean_training_2<- colnames(training[, -58])
testing<- testing[clean_training]
testing_2<- testing_2[clean_training_2]
```

```
## Error in eval(expr, envir, enclos): object 'testing_2' not found
```

```
dim(testing)
```

```
## [1] 7846 58
```

Lastly, to ensure proper function of the Decision Tree analysis, we must coerce the data

```
for (i in 1:length(testing) ) {
  for(j in 1:length(training)) {
    if( length( grep(names(training[i]), names(testing)[j]) ) ==1 ) {
      class(testing[j]) <- class(training[i])
    }
  }
}
testing <- rbind(training[2, -58] , testing)
```

```
## Error in rbind(deparse.level, ...): numbers of columns of arguments do not match
testing <- testing[-1,]
```

K-Fold Cross Validation

We will then use the K-Fold process to cross-validate the data by splitting the training set in to many, smaller data sets.

Here, I am creating 10 folds and setting a random number seed of 32323 for the study. Each fold has approximately the same number of samples in it.

```
set.seed(32323)
folds <- createFolds(y=train_data$classe,k=10,list=TRUE,returnTrain=TRUE)
sapply(folds,length)
```

```
## Fold01 Fold02 Fold03 Fold04 Fold05 Fold06 Fold07 Fold08 Fold09 Fold10
## 17660 17660 17661 17660 17659 17658 17660 17660 17660 17660
```

```
folds[[1]][1:10]
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

Here, I wanted to resample the data set

```
set.seed(32323)
folds <- createResample(y=train_data$classe,times=10,
                        list=TRUE)
sapply(folds,length)
```

```
## Resample01 Resample02 Resample03 Resample04 Resample05 Resample06
## 19622 19622 19622 19622 19622 19622
## Resample07 Resample08 Resample09 Resample10
## 19622 19622 19622 19622
```

```
folds[[1]][1:10]
```

```
## [1] 2 3 3 4 6 6 7 8 9 10
```

Machine learning algorithm decisioning

First I wanted to determine the optimal machine learning model to use. The first test I used the Decision Tree approach. I followed up by testing the Random Forest approach.

Machine learning using Decision Trees

The first task was to determine model fit.

```
## Error in eval(expr, envir, enclos): could not find function "rpart"
## Error in eval(expr, envir, enclos): object 'modelFit' not found
```

Next, construct a Decision Tree graph

```
fancyRpartPlot(modelFit)
```

```
## Error in fancyRpartPlot(modelFit): object 'modelFit' not found
```

Next, I used the predict function for model fitting.

```
predict_perf<- predict(modelFit, testing, type = "class")
```

```
## Error in predict(modelFit, testing, type = "class"): object 'modelFit' not found
```

Lastly, I used a confusion matrix to test the results.

```
confusionMatrix(predict_perf, testing$classe)
```

```
## Error in confusionMatrix(predict_perf, testing$classe): object 'predict_perf' not found
```

The output for this approach is not too bad. Overall model accuracy was 0.86 or 86%. Within a 95% probability, the model accuracy ranges between 0.85 and 0.87. This test is confirmed by a p-value < 0.05.

Machine learning using Random Forests

Like the previous analysis, I wanted to determine model fit. This time instead of using the Decision Tree approach, I used the randomForest function.

```
modelFit2<- randomForest(classe ~., data = training)
```

Next, I wanted to predict the in-sample error.

```
predict_perf2<- predict(modelFit2, testing, type = "class")
```

And once again, the last step was to use a confusion matrix to test results.

```
confusionMatrix(predict_perf2, testing$classe)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    A    B    C    D    E
##           A 2231    0    0    0    0
##           B    0 1518    3    0    0
##           C    0    0 1364    2    0
##           D    0    0    1 1283    5
##           E    0    0    0    1 1437
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.9985
```

```
##           95% CI : (0.9973, 0.9992)
```

```
##           No Information Rate : 0.2844
```

```
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.9981
```

```
##           McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity          1.0000   1.0000   0.9971   0.9977   0.9965
## Specificity          1.0000   0.9995   0.9997   0.9991   0.9998
## Pos Pred Value       1.0000   0.9980   0.9985   0.9953   0.9993
## Neg Pred Value       1.0000   1.0000   0.9994   0.9995   0.9992
## Prevalence           0.2844   0.1935   0.1744   0.1639   0.1838
```

## Detection Rate	0.2844	0.1935	0.1739	0.1635	0.1832
## Detection Prevalence	0.2844	0.1939	0.1741	0.1643	0.1833
## Balanced Accuracy	1.0000	0.9998	0.9984	0.9984	0.9982

After running the data through the Decision Tree and the Random Forests frameworks, the output suggests that the Random Forest approach is superior with this data set. Model accuracy is 0.9978 or 99.7% with a 95% probability that the model has an accuracy between 0.9965 and 0.9987. This test is confirmed by a p-value of < 0.05 .

Finally, we will use the Random Forest model for prediction

```
predict_pef_final<- predict(modelFit2, testing, type = "class")
```

```
##
```

```
##
```

```
## processing file: Peterson - Machine Learning Project.Rmd
```

```
## Error in parse_block(g[-1], g[1], params.src): duplicate label 'setup'
```